




(a) CIFAR-10 vs CIFAR-100
64-D C2ST, 200 repeats



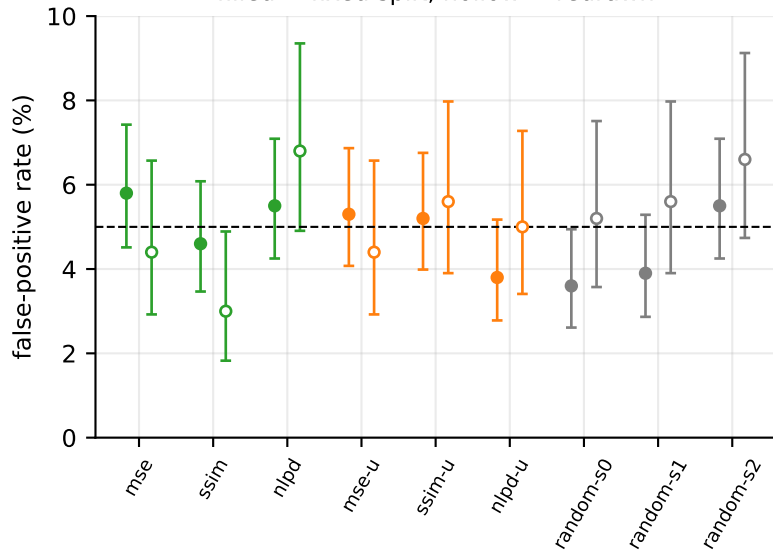
(b) CIFAR-10 vs one class
(detection is 100% everywhere)



(c) drop 1 of 10 classes, $n = 2000$
bar = 64-D, tick = scalar



(d) calibration, 64-D C2ST, $n = 1000$
filled = fixed split, hollow = redrawn



— trained on natural images — untrained (random init)
— trained on uniform noise — JPEG bytes (scalar KS)